

Learned representations, not architecture, drive the gap

